

30563 – MATHEMATICAL MODELLING FOR NEUROSCIENCE
Final Project Report

The Shape of a Stimulus

*Study of geometry of three stimulus classes
in mouse visual cortex (MICrONS)*

Francesca Casillo Alessandro Lombardi
Edoardo Panella Federico Scaffidi Muta

May 13, 2026

1 Introduction

A stimulus, to visual cortex, is not a property of any single neuron. It is a pattern: thousands of cells firing together, occupying a particular region of activity space, and the geometry of that region (where it sits, how it drifts, how far it stands from other stimuli) is what carries class identity [2]. Two of the tools developed in the course probe this geometry from opposite ends. PCA goes looking for the axes that explain the most variance, ignoring any labels we might have. A supervised decoder does the opposite: it goes straight for the directions that tell the classes apart. When the two methods disagree, that disagreement is itself the result: it tells us that the class signal lives in low-variance directions PCA was never asked to find.

The MICrONS Functional Connectomics dataset [1] makes this kind of comparison practical at scale. A single recording session captures roughly 8,000 neurons simultaneously across primary visual cortex (V1) and three higher visual areas (AL, LM, RL), exposed to three stimulus classes that span the natural-to-synthetic range. V1 is where cortical vision starts; AL, LM and RL are its dorsal-stream targets. Comparing the four areas asks a simple question: does class structure sharpen as we move up the hierarchy, or is it already there at the entry point?

2 Research Question

The primary question is straightforward. Do Clip (natural movies), Monet2 (textured noise) and Trippy (synthetic waves) occupy distinguishable population states, and is that separation equally strong across V1, AL, LM and RL? A first round of results forced a second, sharper question. Is what we are seeing a real three-way structure, or just a one-dimensional naturalness boundary in disguise? And if it is a boundary, does any one area specialise on it on a per-neuron basis?

3 Methodology

Session selection. The fourteen functional sessions in MICrONS were screened on five summary statistics (neuron count, median per-neuron activity, mean pupil diameter, mean treadmill speed, and neuron distribution across brain areas), with anything outside $|z| > 2$ flagged as outlier (Fig. 1). Session 7_5 came through clean on every measure: 8,194 neurons against a cohort median of 8,176, no behavioural red flags, and a representative profile across the board.

Preprocessing. Eleven trials with running speed above $|v| > 1$ were dropped, 2.4% of the session, but enough to dominate the response-treadmill correlation we found during EDA. Each neuron’s full timeseries was then linearly detrended (removing the $\sim 45\%$ photobleaching drift identified earlier) and z-scored. 453 trials remained: 377 Clip, 38 Monet2, 38 Trippy. The trial-averaged population response after this pipeline already separates the three stimulus classes by mean activity in every cortical area (Fig. 2).

Three analyses, building on each other. Experiment 1 collapses each trial to a single mean activity vector, fits PCA per area, and quantifies separability two ways: balanced silhouette in the top three PCs (label-agnostic), and 5-fold cross-validated multinomial logistic regression on the full neuron matrix (supervised). Experiment 2 keeps the time axis. For each area and class it builds a trial-averaged $75 \times N$ trajectory and measures pairwise Euclidean distance frame by frame, in two spaces: the top-3 PC subspace (the picture you see in figures) and the full neuron space (the test we trust). Significance is assessed by 100 label-permutation shuffles, variability by 1,000 class-stratified bootstrap resamples, with the PCA basis held fixed throughout, so what the envelopes and nulls capture is uncertainty in the trajectories themselves, not in the projection. Experiment 3 runs the same trajectory pipeline inside Clip, using the three content sub-categories MICrONS provides (Cinematic, sports1m, Rendered).

Controls. Two confounds were obvious. Clip has roughly ten times the trials of either synthetic class, so we sub-sampled it down to 38 and recomputed. V1 has more than thirteen times AL’s

neuron count, so every area was sub-sampled to AL’s 414 neurons (20 random draws) and the entire pipeline rerun.

4 Results

4.1 The signal exists; PCA is not aligned with it (Exp. 1). Each trial reduces to a single N -dimensional vector: component i holds neuron i ’s mean response averaged over the trial’s time window. The session becomes a (trials $\times$ neurons) point cloud, and both the silhouette and the classifier see the same cloud. Cross-validated logistic regression on it reaches 98.9% in V1, 98.0% in LM, 97.8% in RL and 94.7% in AL, comfortably above the 83.2% majority baseline. Viewed through the top three principal components of that same cloud the picture changes: balanced silhouette hovers near zero in every area (Fig. 3). The class signal is there, it is decodable, but it does not live in the directions PCA prioritises. A standard PC-space scatter would have shown nothing (Fig. 4), which is what pushed us toward the time-resolved trajectory analysis that follows.

4.2 Trajectories: an artefact, then the real answer (Exp. 2). Now we retain the time axis. For each area and class, same-class trials are averaged frame by frame to give a trajectory: 75 successive points in neuron space, one per frame, each an N -dimensional vector holding the across-trial mean activity at that timepoint. Monet2 and Trippy trials run for 113 frames natively; we truncated them to Clip’s 75 so the three trajectories are directly comparable. Three classes per area, three trajectories. Plotted in the PC1–PC2 plane, every area produces three clean, stationary, non-overlapping clusters (Fig. 5); stationary because the 75 frames bunch together rather than tracing a path, so each class settles at a fixed address in state space. A first significance test, run in that same top-3 PC space, returned every pair separated at $p \approx 0.01$. But that test cheats. The PCA basis was fit on real-label trial averages, and the shuffled trajectories were then projected through the same basis, so the projection itself carries label information and the shuffle null can never reach what the real data reach. Re-running the test in the full neuron space, where the projection plays no role, turns the picture inside out (Fig. 6). Table 1 reports the honest-test p -values. Monet2 $\leftrightarrow$ Trippy does not separate anywhere. The only synthetic pair that survives is Clip $\leftrightarrow$ Trippy in RL ($p = 0.02$). The Clip-vs-synthetic distinction, on the other hand, holds up across every area. Both controls behave as expected: Clip sub-sampled to 38 trials reproduces the same curves, and matching every area to AL’s neuron count collapses V1’s apparent advantage, leaving the four areas in a tight per-neuron band (Fig. 7).

4.3 Inside Clip: AL as a naturalness detector (Exp. 3). Applying the same trajectory pipeline to Clip’s three content sub-categories, every pair separates at $p \leq 0.02$ in every area, with onset by frame 2 (~ 267 ms post stimulus onset, Fig. 9). At the 7.5 Hz sampling rate, no V1-first hierarchy is resolvable. The interesting result shows up after population matching. AL outperforms every other area on the two pairs that cross the natural/synthetic boundary, Cinematic $\leftrightarrow$ Rendered and sports1m $\leftrightarrow$ Rendered, but matches the others on Cinematic $\leftrightarrow$ sports1m, which is natural-vs-natural (Fig. 10). The simplest reading is that AL is not encoding content category at all. It is encoding the same naturalness axis Experiment 2 turned up, now recurring within Clip itself.

5 Discussion

The answer to the primary question is partly yes. The Clip-vs-synthetic boundary is robust to every control we ran and shows up in every area, so cortex does carry that signal linearly and uniformly across the visual hierarchy. The finer Monet2-vs-Trippy distinction is a different story: it is invisible to linear methods at the population scale. PCA can be made to display it, but only because the basis itself was informed by the labels. Both experiments end up pointing at the same organising dimension. Naturalness explains the inter-class structure in Experiment 2 and reappears as the axis along which AL gains a per-neuron edge in Experiment 3, which together suggest functional differentiation along a natural-vs-artificial axis rather than along content category.

Fitting PCA on real labels and then applying that same projection to shuffled data is a self-confirming test: as soon as the projection step is informed by the labels, the null is compromised, and significance has to be assessed in the unprojected space. We learned this during Experiment 2, which is why the same protection was built into Experiment 3 from the start.

Four limitations qualify the biological reading. *(i)* We cannot cleanly separate category encoding from low-level image statistics (brightness, contrast, motion energy) without stimuli explicitly matched on those statistics, which the dataset does not provide. *(ii)* Everything reported here comes from a single session; cross-mouse generalisation is untested. *(iii)* Linear methods may simply be missing non-linear structure: the Monet2 $\leftrightarrow$ Trippy collapse could be a limitation of the tool rather than of the code. *(iv)* We did not touch the connectomics half of MICrONS, so the structural-to-functional link is still open.

Three things would push this forward. First, a non-linear contrastive method such as CEBRA [3], aimed squarely at the Monet2 $\leftrightarrow$ Trippy collapse, to see whether something non-linear finds what PCA misses. Second, cross-session replication of AL’s apparent naturalness-detector role, to check it is not a quirk of session 7_5. Third, a test of whether AL’s per-neuron advantage aligns with the wiring patterns exposed by the EM volume, which is the question MICrONS was built to make answerable.

References

- [1] The MICrONS Consortium. *Functional connectomics spanning multiple areas of mouse visual cortex*. Nature, 2025.
- [2] P. Dayan and L. F. Abbott. *Theoretical Neuroscience*. MIT Press, 2001.
- [3] S. Schneider, J. H. Lee, M. W. Mathis. *Learnable latent embeddings for joint behavioural and neural analysis*. Nature, 617, 360–368 (2023).

A Tables

Area	Clip $\leftrightarrow$ Monet2	Clip $\leftrightarrow$ Trippy	Monet2 $\leftrightarrow$ Trippy
V1	0.58	0.06	0.37
AL	0.45	0.22	0.76
LM	0.29	0.09	0.48
RL	0.10	0.02	0.29

Table 1: Honest-test p -values: maximum full-feature trajectory distance against the 100-shuffle null. Monet2 $\leftrightarrow$ Trippy never crosses the threshold in any area.

A Figures




Figure 1: **Cohort screening across the 14 MICrONS sessions.** Per-session maximum $|z|$ across five summary statistics (neuron count, median per-neuron activity, mean pupil diameter, mean treadmill speed, fraction of unknown trials). Coral bars cross the $|z| = 2$ flag; session 7_5 (white bar) sits comfortably below the threshold on every axis and was selected for the deep dive.




Figure 2: **Trial-averaged PSTH per stimulus class, per cortical area.** Mean population response averaged within each stimulus class and plotted against trial time, one panel per area (V1, AL, LM, RL). Trippy sits highest, Clip intermediate (with an onset transient in the first few frames), Monet2 lowest. Class identity is already encoded in the first moment of population activity, before any dimensionality reduction.




Figure 3: **Two separability metrics across cortical areas (Exp. 1).** Left: balanced silhouette in the top three PCs, with per-area shuffle-null IQR overlaid as a grey vertical bar. Right: 5-fold cross-validated multinomial logistic regression accuracy on the full neuron matrix, with the 83.2% majority baseline as the dashed reference. The two metrics disagree: silhouette hovers around zero in every area, classifier accuracy is at ceiling. The class signal is decodable but is not aligned with the high-variance axes PCA picks.




Figure 4: **Trial-averaged PCA per area in PC1-PC2 (Exp. 1).** One panel per cortical area (V1 top-left, AL top-right, LM bottom-left, RL bottom-right). Coloured dots are individual trials; crosses mark class centroids. The three stimulus classes overlap heavily in the high-variance subspace in every area: PCA finds the directions of largest variance, but those directions are not the ones that separate the classes.




Figure 5: **Trial-averaged trajectories in PC1–PC2 space, per area (Exp. 2).** Four panels (V1, AL, LM, RL): three stationary, non-overlapping clusters per area. The filled circle marks frame 0; the open circle marks the last frame. The frame-0 markers sit offset from each cluster’s bulk, indicating a shared neutral starting point and a rapid jump to stimulus-specific regions of state space within the first few frames (the onset transient seen in Fig. 2).




Figure 6: **V1 pairwise trajectory distance time courses: biased vs honest (Exp. 2).** Left panel: distances computed in the full neuron space (the honest test). Right panel: distances computed in the top-3 PC subspace fit on real-label trial averages (the biased test). Solid lines are observed pairwise distances per frame; shaded bands are bootstrap envelopes (B=1000); dashed horizontals are the 95th percentile of the 100-shuffle null max distance per pair. The biased test puts every pair above its null. The honest test does not: Monet2↔Trippy stays inside its null band throughout the trial.




Figure 7: **Population-size control on Monet2↔Trippy (Exp. 2)**. Cross-area Monet2↔Trippy distance time course after subsampling V1, LM, RL to AL’s 414 neurons (20 random draws, median $\pm$ IQR shown). Left: full-feature metric. Right: top-3 PC metric. V1’s apparent lead in the raw comparison vanishes once neuron counts are equalised; all four areas converge to the same per-neuron band, confirming that V1’s nominal advantage was a population-size effect rather than a per-neuron one.




Figure 8: **Clip-category trajectories in PC1–PC2 space, per area (Exp. 3)**. Same layout as Fig. 5 but coloured by Clip content sub-category (Cinematic, sports1m, Rendered) rather than by stimulus class. All three categories separate cleanly in every area and remain stationary within their own region of state space after the onset transient.




Figure 9: **Cross-area Clip-category distance time courses, all three pairs (Exp. 3).** Full-feature pairwise distance per frame, one panel per category pair (Cinematic↔sports1m, Cinematic↔Rendered, sports1m↔Rendered), with one coloured line per cortical area and bootstrap envelopes shaded. All four areas separate all three pairs at $p \leq 0.02$; onset is at frame 2 (~ 267 ms post stimulus onset).



Figure 10: **Within-Clip sub-category distances after population matching (Exp. 3).** Same three pairs as Fig. 9 but with every area subsampled to AL's 414 neurons (20 random draws, median $\pm$ IQR). Top row: full-feature metric. Bottom row: top-3 PC metric. AL leads on Cinematic↔Rendered and sports1m↔Rendered (the two natural-vs-synthetic contrasts) but collapses into the other areas on Cinematic↔sports1m (natural-vs-natural). The naturalness axis from Exp. 2 recurs within Clip.